

P R O J E C T :



RATED
PJ

FOR: POD JOBS
DYNAMIC MULTI-AGENT SWARMS
UNDER HUMAN SUPERVISION

KAGGLE 5-DAY INTENSIVE: AI AGENTS CAPSTONE PROJECT

June 15 - 19, 2026 • Hosted by Google



Architect & Author: Devs One — Danny Bouldiez

Platform: Google AI Studio, Google Antigravity, and Nous Research Hermes Agent

Build Status: Deployed & Live

On-Chain	Attestation	Video	NFT:	[Tezos	Objkt	Token
#94](https://objkt.com/tokens/KT1L2gY2BUE2gcydLUXLzSAYwAvriYvZMBQ8/94)						

PROJECT PJ: A COGNITIVE AMPLIFIER FOR PROFESSIONAL SWARMS

■ Executive Summary

Conventional enterprise workflows scale horizontally by adding headcount—introducing communication overhead, scheduling delays, and structural coordination errors. **Project PJ (PodJobs.ai)** proves a new paradigm of **Cognitive Amplification**: *one human specialist directing a synchronized, parallel consensus department of 12 intelligent AI agents*.

By combining the **Google GenAI SDK**, **Model Context Protocol (MCP)**, **NVIDIA NeMo Guardrails emulation**, and **cryptographic Merkle Root consensus signatures**, Project PJ provides a secure, zero-overhead workflow sandbox that scales capabilities instead of headcount.



GitHub: [DannyB-bit/PodJobs](#)



Live Website: [podjobs.vercel.app](#)



On-Chain Walkthrough NFT

Disclaimer: The characters, organizations, and events depicted in the creative components of this document are fictional. Any resemblance to actual persons, living or dead, or real-world entities is purely coincidental and not intended by the author.

■ The On-Going PLOT and The Solution

The Corporate Harvest: When Job Hunting Becomes a Data Mine

Imagine a future where you've spent six years mastering blockchain, Web3, neural networks, and the complex plumbing of edge IoT. You are ready to build the future. But the job market you step into isn't looking to hire you—it's looking to *harvest* you.

Today, U.S. job seekers face a predatory trend. You submit your resume, jump through hoops, and get called into multi-stage technical interviews. But these interviews are no longer evaluation channels. Behind the scenes, managers deploy AI systems to extract your database schemas, scrape your deployment techniques, and farm your architectural knowledge. Once your mind is cataloged into their training models, their system auto-fires a generic, copy-paste template rejection: *"We chose a candidate who matches our needs more closely."* No pay. No compensation. Your intellectual property has been farmed, and you are left facing financial hardship.

The Rise of AI-Driven Restructuring

This isn't an isolated headache; it is a systemic crisis. According to data from Challenger, Gray & Christmas, U.S. employers announced a massive surge of **38,579 AI-driven job cuts in May 2026 alone**—accounting for nearly 40% of all layoffs that month. We see giants like Block cutting 4,000 roles to make way for their proprietary "BIRD" automation engine, and Coinbase slashing 14% of their staff to force engineers into "single-person pods."

The corporate theology is simple: *maximize profits, minimize expenses, displace the human*. But turning families' lives upside down just to pad quarterly balance sheets is not innovation. It is a design failure.

The PJ Vision: Reclaiming Autonomy with Recycled E-Waste

We refused to be automated out. We refused to let our knowledge be farmed.

So, we gathered the e-waste collecting dust in our closets—old Raspberry Pi boards and Helium RAK miners from the early Web3 mining boom. We flashed them, sandboxed them, and turned them into **Project PJ (PodJobs.ai)**.

Project PJ is not an app or a digital commodity. It is a paradigm shift. Instead of displacing workers, we teach them to pilot their own **Sovereign Agent Pods**. We transition the human from a vulnerable cog to a high-productivity **Conductor** of a localized, 12-agent consensus department. By utilizing local offline execution frameworks like **Nous Research Hermes** and running them at the edge, the data never leaves your room.

The corporate solution is fear and layoffs. The **Project PJ** solution is remote work people love, local data sovereignty, and human-amplified productivity. We aren't replacing the team. We are giving the worker the power of a department.

■ ■ The Challenge & The Solution

The Corporate Problem	Project PJ Solution
Execution Latency: 3-10 days for complex documents.	Asynchronous Swarms: Multi-agent tasks finished in under 15 seconds.
Verification Overhead: Peer-reviews introduce delays.	Bayesian Consensus: Unanimous 12-node voting verification.
Security Risks: Confidentiality leaks and input injection.	Isolation Shield: Prompt sanitization and NeMo Guardrail thresholds.
Audit Trails: Hard to track steps in multi-agent reasoning.	Merkle-Attestation Seals: Immutable cryptographic signatures.

■■ Architecture & Sequential Cascade (ADK)

Project PJ implements a 5-node sequential execution cascade utilizing the **Google GenAI SDK (gemini-3.5-flash)**:

```
graph TD
  A[Human Director Prompt] --> B[Planner Node]
  B -->|Deconstructs goals| C[Context Miner RAG Node]
  C -->|Mines case files / documents| D[Suite Draftsman Node]
  D -->|Synthesizes professional briefs| E[System Integrity Auditor Node]
  E -->|Scans for safety & biases| F[Swarms Consensus Arbiter Node]
  F -->|Bayesian consensus seal| G[Merkle Root Cryptographic Seal]
```

1. **Planner Node**: Maps out execution branches and allocates agent responsibilities.
2. **Context Miner Node**: Performs vector grounding (RAG) based on local knowledge databases.
3. **Suite Draftsman Node**: Compiles professional documents and templates.
4. **System Integrity Auditor Node**: Emulates **NVIDIA NeMo Guardrails**, rejecting any payload with a bias/safety score exceeding **0.05**.
5. **Swarms Consensus Arbiter Node**: Performs final validation and returns the output along with a mathematical seal.

Cryptographic Consensus & Merkle Attestation

To guarantee the logical integrity of the reasoning path, the input, actions, and outputs of all agent nodes are sequentially hashed and sealed into a **Merkle Tree**. The resulting **Merkle Root Hash** acts as an immutable tamper-proof signature, proving that no node's intermediate outputs were modified during the cascade.

■ Local CLI & MCP Extensibility

To enable external developer tools and shell automation, Project PJ exposes:

- **JSON-RPC MCP Server (mcp-server/index.js)**: A stdio-based server compliant with the Model Context Protocol, exposing swarm simulation and inspect tools to IDEs like Claude Desktop.
- **Command-Line Agent Skill (bin/podjobs-cli.js)**: A CLI enabling offline simulations and live sequential runs directly from the developer shell.
- **Deployed API Validator (bin/validate-live-api.js)**: A testing harness verifying that Vercel serverless routes correctly connect with live Google APIs.

■ The Nous Research Hermes Connection

Project PJ integrates the state-of-the-art **Nous Research Hermes Agent** configuration framework. By utilizing Nous Research blueprints, each node loaded in the dynamic workstation utilizes customized **hermes.json** onboarding manifests, configuring:

- Strategic organizational directives.
- Persona alignment parameters.
- Isolated model instruction sets.

License Credit & Attestation

We extend our deepest gratitude to the **Nous Research** team for their contributions to open-source agent architectures. Nous Research Hermes and its derivative configurations are utilized in this project under the terms of the Apache License 2.0.

```
Apache License
Version 2.0, January 2004
http://www.apache.org/licenses/

Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.
```

■ Robust Fallback & Live Stress Test

When utilizing commercial LLM API integrations, operational stability is a key challenge. If the system is subjected to rate limiting or connection failure, Google does not automatically switch models on their side; they will simply return an HTTP 429 (Too Many Requests) or a server error.

To protect the user experience and ensure 100% service uptime, Project PJ incorporates a custom-built, self-healing try-catch fallback router within [app/api/gemini/route.ts](file:///k:/Cpastone-Project-kaggle5day/app/api/gemini/route.ts).

How the Fallback Router Works:

- 1. **Live Attempt:** The system attempts to run the dynamic multi-agent cascade using the active **gemini-3.5-flash** model.
- 2. **API Interception:** If Google returns a 429 Rate Limit, quota error, or any other API connection failure, the try-catch block intercepts the error.
- 3. **Instant Mitigation:** The backend immediately switches the execution engine to the **Fallback Simulation Engine (Live Error)**. It generates realistic, prompt-aligned agent outputs and calculates a valid cryptographic Merkle Root, returning a successful response.
- 4. **Zero UI Disruption:** Because of this self-healing script, the frontend dashboard, CLI outputs, and Merkle Tree animations never crash or freeze.

Live Stress Test Attestation

To verify this behavior, we programmatically stress-tested the live production endpoints at podjobs.vercel.app under concurrent load. The results are summarized below:

```
=====
STRESS TEST SUMMARY REPORT
=====
Phase 1 (Live Run):
- Concurrency duration: 6.89s
- Handled by Live Gemini 3.5 Flash Engine: 15/15 (100%)

Phase 2 (Forced Failure Run - Invalid Key simulation):
- Concurrency duration: 0.55s
- Safely Intercepted & Handled by Fallback Engine: 15/15 (100%)

Overall System Resilience:
- Total API requests sent: 30
- Total successful HTTP 200 responses: 30/30
- Uptime Success Rate: 100.0%
✓ PASS: Fallback router and live endpoints are fully resilient and self-healing!
```

What this proves:


- **Live Concurrency:** The live production server successfully executed 15 concurrent workflows (totaling 75 sequential live Gemini API calls) in 6.89 seconds under standard conditions.
- **Graceful Failure:** Under simulated API failures (Phase 2), the system caught 100% of errors and triggered the fallback simulation engine in under 0.55 seconds.
- **Absolute Uptime:** Zero requests failed, confirming complete robustness.

■ Attestation & Sign-off

"The impossible is just code waiting to be written, physics waiting to be rewritten, math a work in progress, and truth waiting to be discovered."

I AM

— **Devs One** (Danny Bouldiez)




HUMAN - DNA SEQUENCE ATTESTATION

Special thanks to Kaggle and Google for hosting the Kaggle 5-Day Intensive AI Agents Vibe Coding Course (2026).

This whitepaper and project are intended for educational and entertainment purposes only. If we offended anyone, go talk to Chat.

The Google logo is displayed in its standard multi-colored font, with the letters 'G', 'o', 'o', 'g', 'l', and 'e' in blue, red, yellow, blue, green, and red respectively.The Kaggle logo is displayed in a solid blue, lowercase, sans-serif font.

WE ARE



Q TheAiCollective.art